



03. Breakng Off



Two-Stage Stochastic Optimization

If DFL targets one-stage stochastic optimization, could we do **two-stage**?



- For example, in **first stage** we decide what to pack in our suitcase
- ...During the trip, we may realize we have **forgotten something**
- ...And we need to **spend money** to buy the missing stuff



Two-Stage Stochastic Optimization

If DFL targets one-stage stochastic optimization, could we do two-stage?

Two-stage problems are among the most interesting in stochastic optimization

- They involve making a set of decisions now
- Then observing how uncertainty unfolds
- ...And making a second set of decisions

The former are called first-stage decisions, the latter recourse actions

Here's an example we will use for this topic

Say we need to secure a supply of resources

- First, we make contracts with primary suppliers to minimize costs
- If there are unexpected setbacks (e.g. insufficient yields)
- ...Then we can buy what we lack from another source, but at a higher cost



Two-Stage Stochastic Optimization

Let's define two-stage stochastic optimization problems (2s-SOP) formally:

$$\operatorname{argmin}_z \left\{ f(z) + \mathbb{E}_{y \sim P(Y|x)} \left[\min_{z''} r(z'', z, y) \right] \mid z \in F, z'' \in F''(z, y) \right\}$$

- Y represents the uncertain information
- z is the vector of **first stage decisions**
- F is the feasible space for the first stage
- z'' is the vector of **recourse actions**
- z'' is not fixed: it can change for every sampled y
- The set of feasible recourse actions $F''(z, y)$ also changes for every y
- f is the immediate cost function, r is the cost of the recourse actions



A Simple Example

We will consider this simple problem

...Which is based on our previous supply planning example:

$$\operatorname{argmin}_z c^T z + \mathbb{E}_{y \sim P(Y|x)} \left[\min_{z''} c'' z'' \right]$$

subject to: $y^T z + z'' \geq y_{min}$
 $z \in \{0, 1\}^n, z'' \in \mathbb{N}_0$

- $z_j = 1$ iff we choose then h -th supply contract
- c_j is the cost of the j -th contract
- y_j is the yield of the j -th contract, which is uncertain
- y_{min} is the minimum total yield, which is known
- z'' is the number of units we buy at cost c'' to satisfy the yield requirement



Scenario Based Approach

Classical solution approaches for 2s-SOP are scenario based

We start by sampling a finite set of N values from $P(Y \mid x)$

$$\begin{aligned} & \operatorname{argmin}_z \min_{z''} c^T z + \frac{1}{N} c'' z''_k \\ & \text{subject to: } y^T z + z''_k \geq y_{min} \quad \forall k = 1..N \\ & z \in \{0, 1\}^n \\ & z''_k \in \mathbb{N}_0 \quad \forall k = 1..N \end{aligned}$$

Then we build different recourse action variables for each scenario

- ...We define the feasible sets via constraints
- ...And we use the Sample Average Approximation to estimate the expectation

The method is effective, but also computationally expensive



DFL for 2s-SOP

Could we tackle 2s-SOP with DFL?

As a recap, our DFL training problem is:

$$\theta^* = \operatorname{argmin}_{\theta} \left\{ \mathbb{E}_{(x,y) \sim P(X,Y)} [\text{regret}(\hat{y}, y)] \mid \hat{y} = h(x, \theta) \right\}$$

With:

$$\text{regret}(\hat{y}, y) = y^T z^*(\hat{y}) - y^T z^*(y)$$

And:

$$z^*(y) = \operatorname{argmin}_z \{ y^T z \mid z \in F \}$$



DFL for 2s-SOP

With the same transformations used in the one-stage case, we get:

$$\theta^* = \operatorname{argmin}_{\theta} \left\{ \mathbb{E}_{y \sim P(Y|X=x)} [y^T z^*(\hat{y})] \mid \hat{y} = h(x, \theta), z^*(\hat{y}) \in F \right\}$$

Now, say we had a DLF approach that could deal with any function $g(z, y)$

- In this case y would be a vector of uncertain parameters (not necessarily costs)
- The function should compute the equivalent of $y^T z^*(\hat{y})$
- ...i.e. the true cost of the solution computed for the estimate costs

Under this conditions, at training time we could solve:

$$\theta^* = \operatorname{argmin}_{\theta} \left\{ \mathbb{E}_{y \sim P(Y|X=x)} [g(z^*(\hat{y}), y)] \mid \hat{y} = h(x, \theta), z^*(\hat{y}) \in F \right\}$$

It would still be DFL, just a bit more general



DFL for 2s-SOP

At this point, let's choose:

$$g(z, y) = \min_{z''} \{ f(z) + r(z'', z, y) \mid z'' \in F''(z, y) \}$$

- For a given solution z , $g(z, y)$ computes the best possible objective
- ...Assuming that the value of the parameters is y

By substituting in the training formulation we get:

$$\operatorname{argmin}_{\theta} f(z^*(\hat{y})) + \mathbb{E}_{y \sim P(Y|x)} \left[\min_{z''} r(z'', z^*(\hat{y}), y) \right]$$

subject to: $\hat{y} = h(x, \theta)$, $z^*(\hat{y}) \in F$, $z'' \in F''(z, y)$



..Which can definitely be used for 2s-SOP problems!

Grouding the Approach

We can ground the approach by relying on the scenario-based formulation

In our example problem, we compute $z^*(y)$ by solving:

$$\begin{aligned} z^*(y) &= \operatorname{argmin}_z \min_{z''} c^T z + c'' z''_k \\ \text{subject to: } &y^T z + z''_k \geq y_{min} \\ &z \in \{0, 1\}^n \\ &z''_k \in \mathbb{N}_0 \end{aligned}$$

And we define $g(z, y)$ as:

$$\begin{aligned} g(z, y) &= \min_{z''} c^T z + c'' z''_k \\ \text{subject to: } &y^T z + z''_k \geq y_{min} \\ &z''_k \in \mathbb{N}_0 \end{aligned}$$



Overview and Properties

Intuitively, the approach works as follows

- We observe x and we compute \hat{y}
 - We compute $z^*(\hat{y})$ by solving a scenario problem
 - We compute $g(z^*(\hat{y}), y)$ by solving a scenario problem with fixed z values
- ...And we end up minimizing the expected cost of the 2s-SOP

We have 1 restriction and 3 "superpowers" w.r.t. the classical approach

- The restriction: we control z^* only through θ
- Superpower 1: we are not restricted to a single x
- Superpower 2: works with any distribution
- Superpower 3: at inference time, we always consider a single scenario



Scalable Two-stage Stochastic Optimization

The last advantage is massive

The weakest point of classical 2s-SOP approach is scalability

- Multiple scenarios are required to obtain good results
- ...But they also add more variables

With NP-hard problem, that solution time may grow exponentially

With this approach, the computational cost is all at training time

- It can even be lower, since you always deal with single scenarios
- There are alternatives, such as [1], where ML is used to estimate the recourse
- ...These have their own pros and cons

[1] Dumouchelle, Justin, et al. "Neur2sp: Neural two-stage stochastic programming." arXiv preprint arXiv:2205.12006 (2022).



The Elephant in the Room

So far, so good, but how to we make $g(z, y)$ differentiable?

There are a few alternatives, all with limitations:

- The approach from [1] handles parameters in the problem constraints
 - It is based on the idea of differencing the recourse action
 - ...But it is (mostly) restricted to 1D packing problems
- The approach from [2] can be used for 2s-SOP with a stretch
 - It based on idea of embedding a MILP solver in ML
 - ...But it's semantic does not fully align with 2s-SOP

Here, we will see a different technique

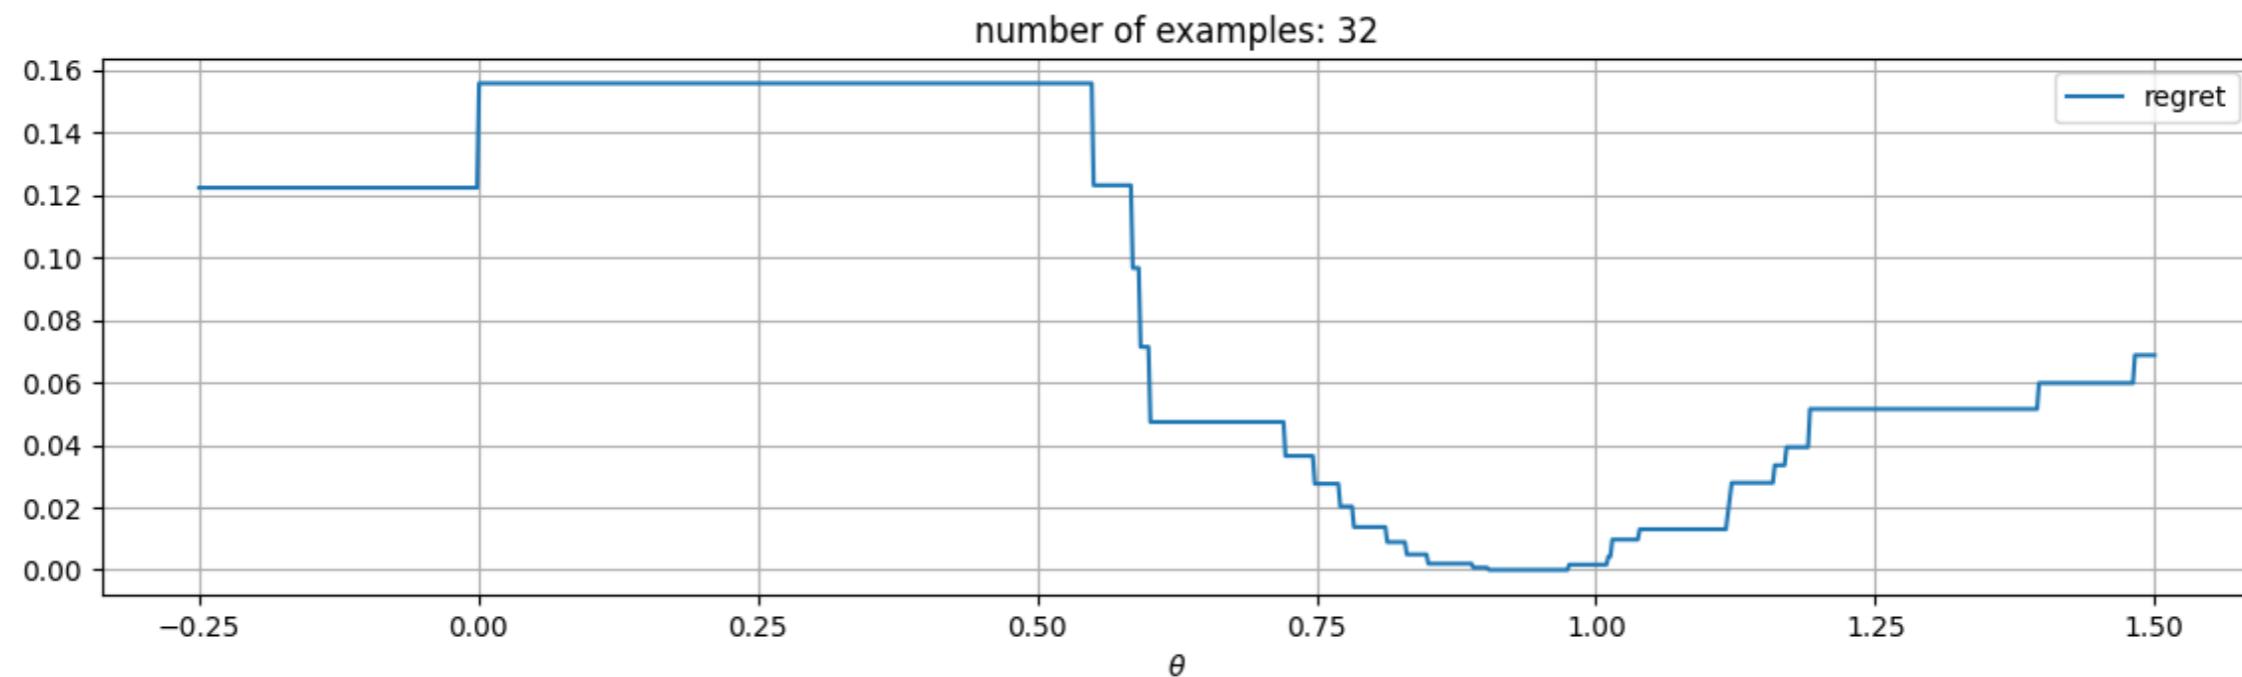
[1] Hu, X., Lee, J. C. H., and Lee, J. H. M. *Predict+optimize for packing and covering lps with unknown parameters in constraints*. CoRR, abs/2209.03668, 2022. doi: 10.48550/arXiv.2209.03668.

[2] Paulis, Anselm, et al. "Comboptnet: Fit the right np-hard problem by learning integer programming constraints." International Conference on Machine Learning. PMLR, 2021.

Looking Back at SPO

Let's look again at the regret loss for our original toy example

```
In [2]: util.draw_loss_landscape(losses=[util.RegretLoss()], model=1, seed=42, batch_size=32, figsize=f)
```



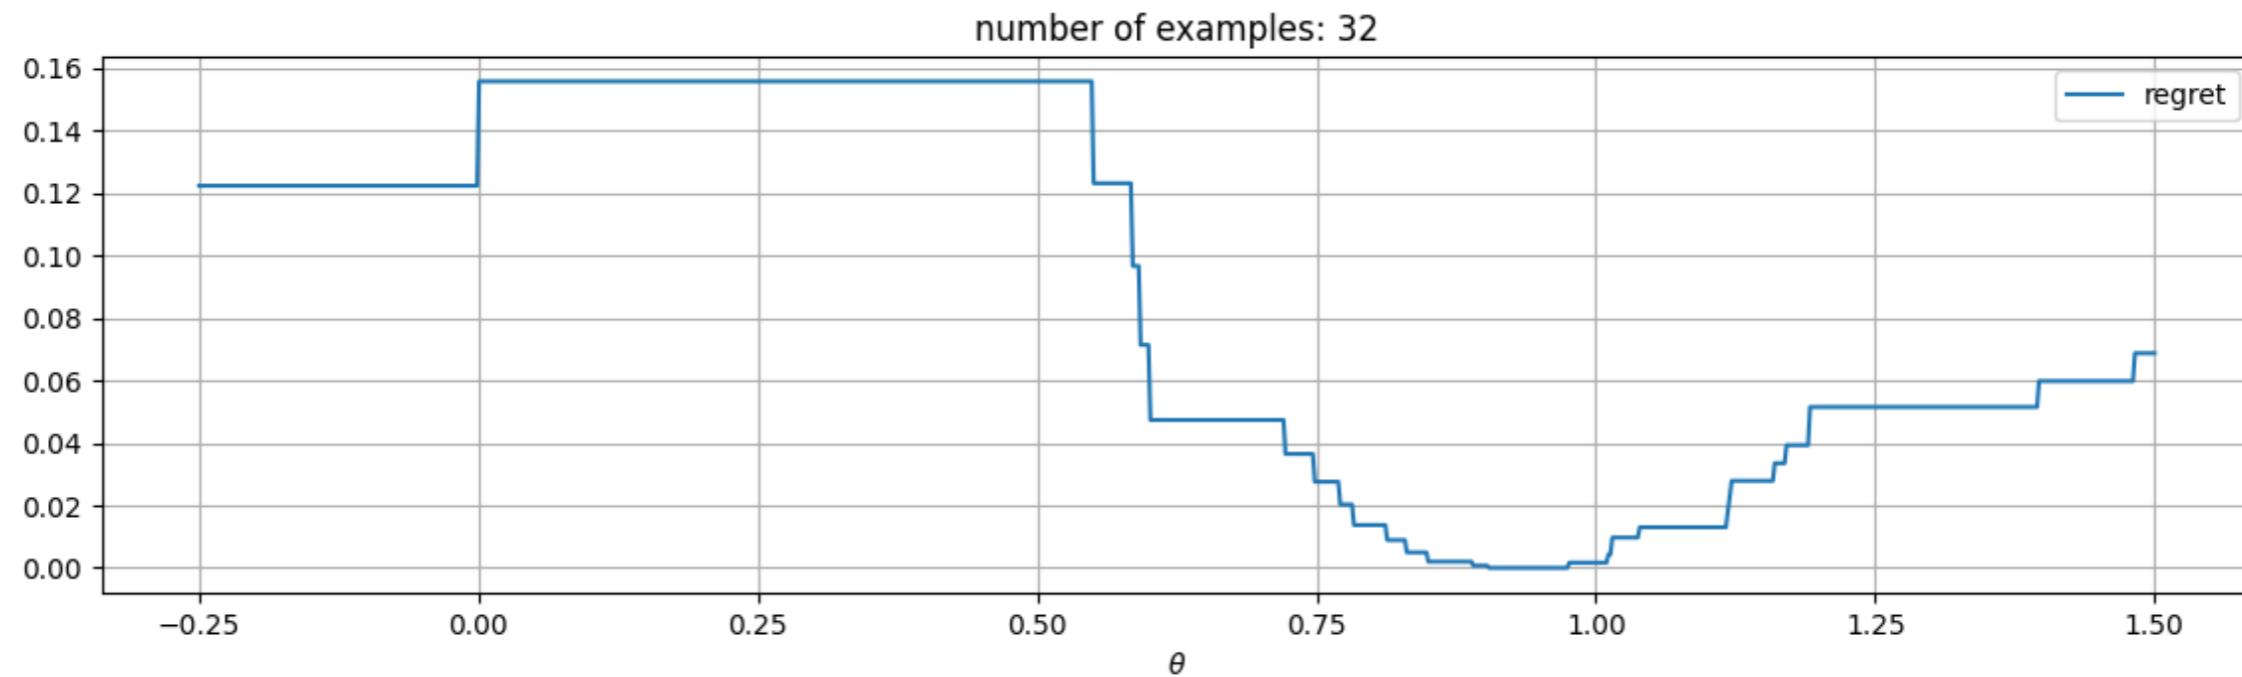
- It is non-differentiable at places, and flat almost everywhere
- Can we think of another way to address these issues?



Looking Back at SPO

If we could act on this function itself, a simple solution would be **smoothing**

```
In [3]: util.draw_loss_landscape(losses=[util.RegretLoss()], model=1, seed=42, batch_size=32, figsize=figsize)
```



- We could think of computing a convolution with a Gaussian kernel
- It would be like applying a Gaussian filter to an image



Stochastic Smoothing

But how can we do it through an optimization problem?

A viable approach is using stochastic smoothing

- Rather than learning a point estimator $h(x, \theta)$
- We learn a **stochastic estimator** s.t. $\hat{y} \sim \mathcal{N}(h(x, \theta), \sigma)$

Intuitively:

- We still use a point estimator, but to predict **a vector of means**
- Then we sample \hat{y} from a normal distribution having the specified mean
- ...And a fixed standard deviation

We end up smoothing over \hat{y} rather than over θ

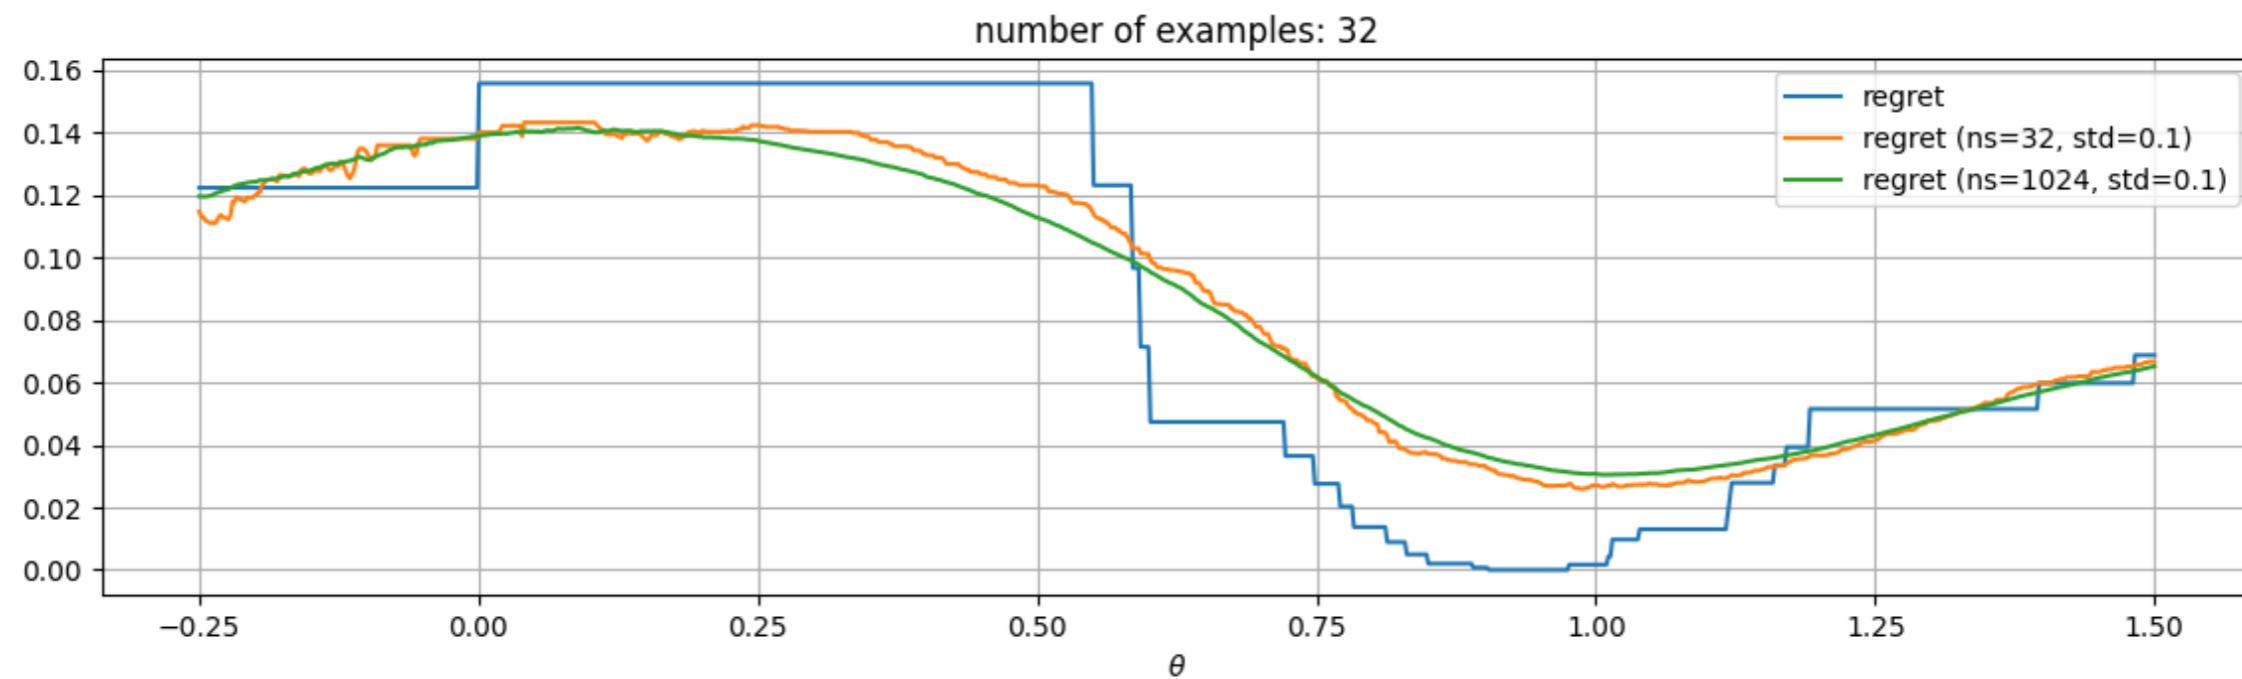
But it's very close to what we wanted to do!



Stochastic Smoothing

Let's see how it works on our toy example

```
In [4]: util.draw_loss_landscape(losses=[util.RegretLoss(), util.RegretLoss(smoothing_samples=32, smooth
```



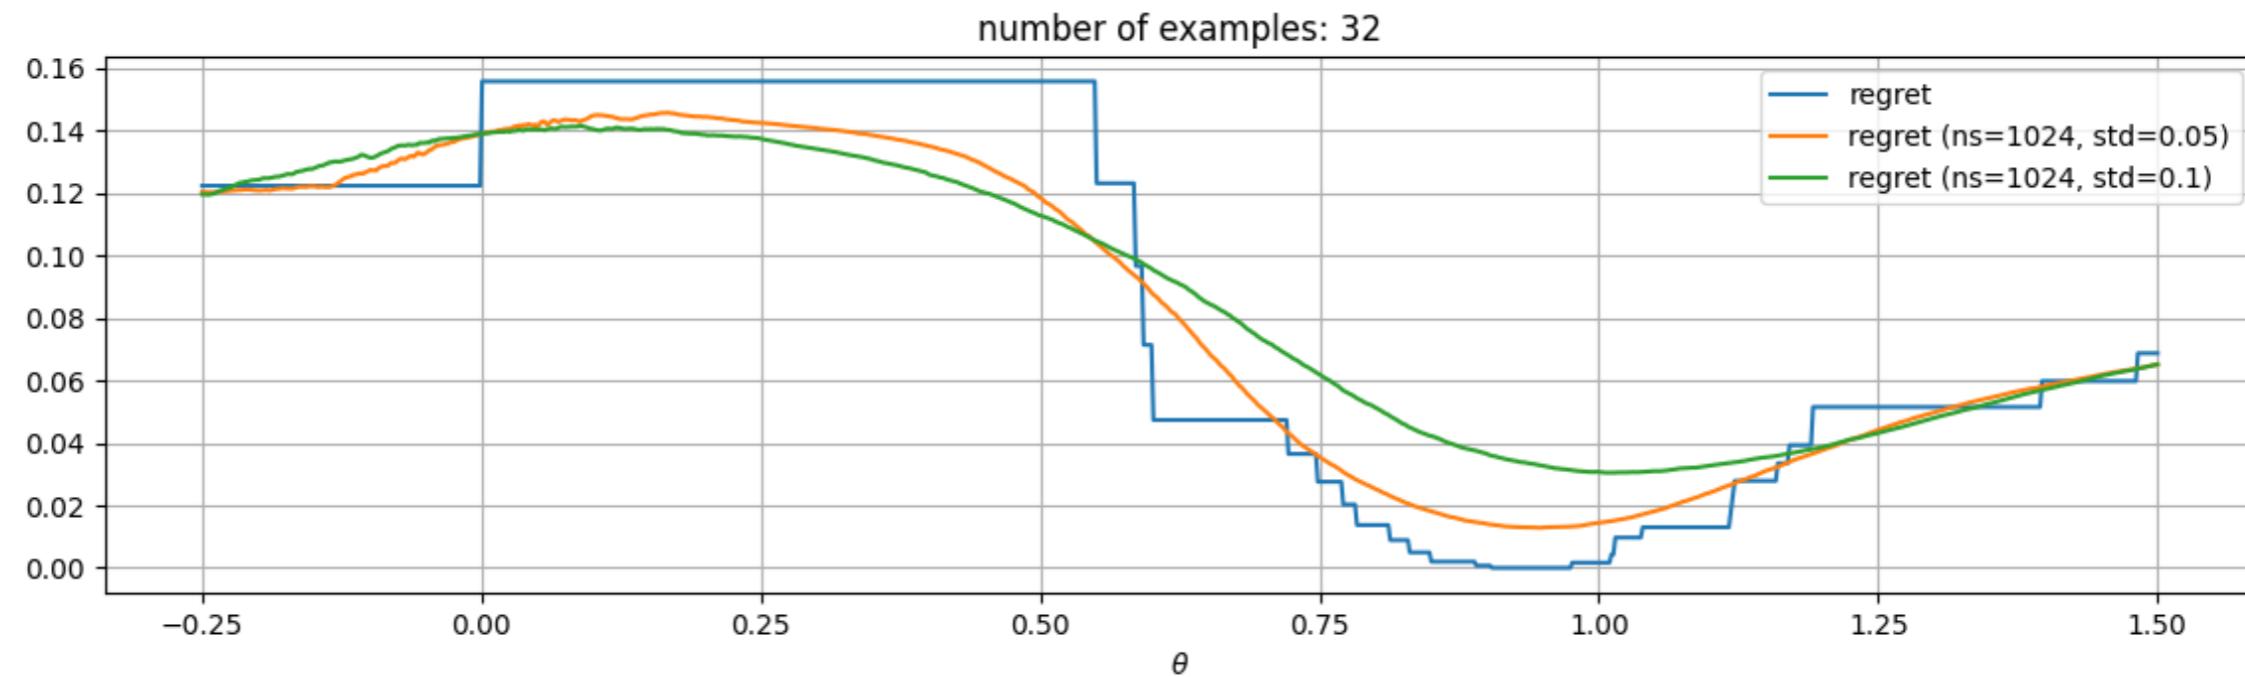
- It's a stochastic approach, some noise is to be expected
- Using more samples leads to better smoothing



Stochastic Smoothing

We can control the smoothing level by adjusting σ

```
In [5]: util.draw_loss_landscape(losses=[util.RegretLoss(), util.RegretLoss(smoothing_samples=1024, smoo
```



- Larger σ value remove flat sections better
- ...But also cause a shift in the position of the optimum



Score Function Gradient Estimation

How does that help us?

Normally, the DFL loss looks like this:

$$L_{DFL}(\theta) = \mathbb{E}_{(x,y) \sim P(X,Y)} [\text{regret}(\hat{y}, y)]$$

When we apply stochastic smoothing, it turns into:

$$\tilde{L}_{DFL}(\theta) = \mathbb{E}_{(x,y) \sim P(X,Y), \hat{y} \sim \mathcal{N}(h(x,\theta))} [\text{regret}(\hat{y}, y)]$$

The expectation is now computed on x , y , and \hat{y}

- We can use a sample average to handle the expectation on x and y
- ...But if we do it on \hat{y} we are left with nothing differentiable



Score Function Gradient Estimation

So we expand the last expectation on \hat{y} :

$$\tilde{L}_{DFL}(\theta) = \mathbb{E}_{(x,y) \sim P(X,Y)} \left[\int_{\hat{y}} \text{regret}(\hat{y}, y) p(\hat{y}, \theta) d\hat{y} \right]$$

- $\text{regret}(\hat{y}, y)$ cannot be differentiated, since \hat{y} is a fixed sample in this setup
- However, the probability $p(\hat{y}, \theta)$ can! It's just a Normal PDF

Now, we just need a way to handle the integral

We do it by **focusing on the gradient**

Due to linearity of expectation and integration, this is given by:

$$\nabla \tilde{L}_{DFL}(\theta) = \mathbb{E}_{(x,y) \sim P(X,Y)} \left[\int_{\hat{y}} \text{regret}(\hat{y}, y) \nabla_{\theta} p(\hat{y}, \theta) d\hat{y} \right]$$



Score Function Gradient Estimation

Let's consider again the expression we have obtained

$$\nabla \tilde{L}_{DFL}(\theta) = \mathbb{E}_{(x,y) \sim P(X,Y)} \left[\int_{\hat{y}} \text{regret}(\hat{y}, y) \nabla_{\theta} p(\hat{y}, \theta) d\hat{y} \right]$$

Since that $\log'(f(x)) = 1/x f'(x)$, we can rewrite the formula as:

$$\nabla \tilde{L}_{DFL}(\theta) = \mathbb{E}_{(x,y) \sim P(X,Y)} \left[\int_{\hat{y}} \text{regret}(\hat{y}, y) p(\hat{y}, \theta) \nabla_{\theta} \log p(\hat{y}, \theta) d\hat{y} \right]$$

Now, the integral is again an expectation, so we have:

$$\nabla \tilde{L}_{DFL}(\theta) = \mathbb{E}_{(x,y) \sim P(X,Y), \hat{y} \sim \mathcal{N}(h(x,\theta), \sigma)} [\text{regret}(\hat{y}, y) \nabla_{\theta} \log p(\hat{y}, \theta)]$$



Score Function Gradient Estimation

Finally, we can use a sample average to approximate both expectations:

$$\nabla \tilde{L}_{DFL}(\theta) \simeq \frac{1}{m} \sum_{i=1}^m \frac{1}{N} \sum_{k=1}^N \text{regret}(\hat{y}, y) \nabla_\theta \log p(\hat{y}, \theta)$$

- For every training example we sample \hat{y} from the stochastic estimator
- We compute $\text{regret}(\hat{y}, y)$ as usual
- ...And we obtain a gradient since $p(\hat{y}, \theta)$ is easily differentiable in θ

We can trick a tensor engine into doing the calculation by using this loss:

$$\tilde{L}_{DFL}(\theta) \simeq \frac{1}{m} \sum_{i=1}^m \frac{1}{N} \sum_{k=1}^N \text{regret}(\hat{y}, y) \log p(\hat{y}, \theta)$$



Score Function Gradient Estimation

This approach is also known as Score Function Gradient Estimation (SFGE)

- It is a known approach (similar to [3]), but it has not been used in DFL
- We applied it to 2s-SOP in [4] (accepted, not yet published)

It works with any function, not just regret

...And in practice it can be improved by standardizing the gradient terms:

$$\nabla \tilde{L}_{DFL}(\theta) \simeq \frac{1}{m} \sum_{i=1}^m \frac{1}{N} \sum_{k=1}^N \frac{g(\hat{y}, y) - \text{mean}(g(\hat{y}, y))}{\text{std}(g(\hat{y}, y))} \nabla \log p(\hat{y}, \theta)$$

- Standardization helps in particular with small numbers of samples

[3] Berthet, Quentin, et al. "Learning with differentiable perturbed optimizers." *Advances in neural information processing systems* 33 (2020): 9508-9519.

[4] Silvestri, Mattia et al. "Score Function Gradient Estimation to Widen the Applicability of Decision-focused Learning", *Differentiable Almost Everywhere workshop at ICML 2023*



A Practical Example

We test this on our supply planning problem

We start by generating a dataset of **contract values** (the costs are fixed)

```
In [21]: seed, nitems = 42, 20  
data_tr = util.generate_costs(nsamples=350, nitems=nitems, seed=seed, noise_scale=.2, noise_type='uniform')  
util.plot_df_cols(data_tr, figsize=figsize, title='Training Set', scatter=True)
```



The distribution is the same we used for the one-stage problem



A Practical Example

Then we generate the remaining problem parameters

```
In [22]: # Generate the problem
rel_req = 0.6
rel_buffer_cost = 10
prb = util.generate_2s_problem(nitems, requirement=rel_req * data_tr.mean().sum(), rel_buffer_cost=rel_buffer_cost)
prb
```

```
Out[22]: ProductionProblem2Stage(costs=[1.14981605 1.38028572 1.29279758 1.23946339 1.06240746 1.062397
81
1.02323344 1.34647046 1.240446 1.28322903 1.0082338 1.38796394
1.33297706 1.08493564 1.07272999 1.0733618 1.1216969 1.20990257
1.17277801 1.11649166], requirement=3.8862101169088654, buffer_cost=11.830809153591137)
```

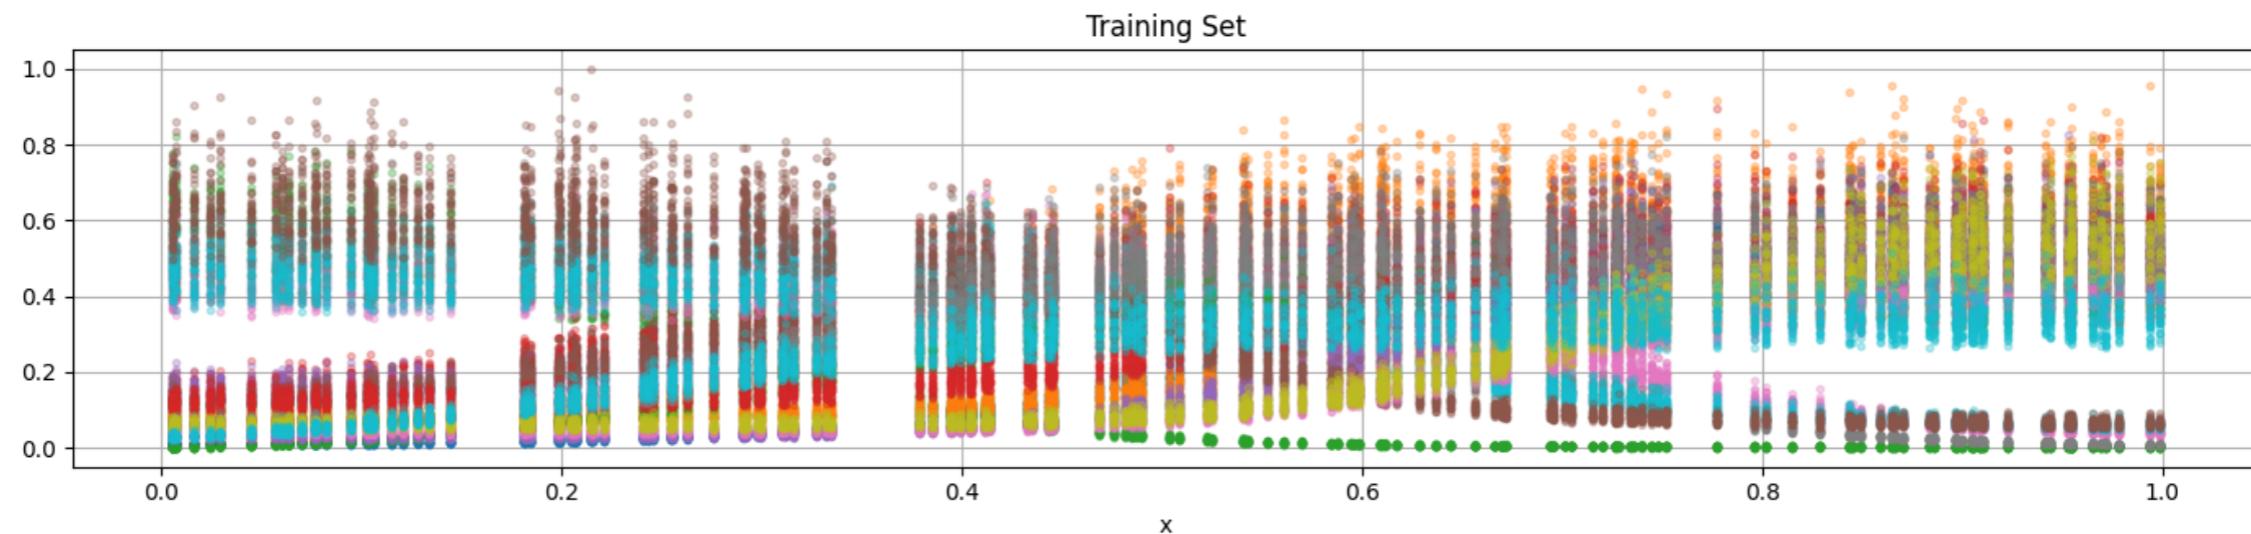
- The minimum value if 60% of the sum of average values on the training data
- Buying in the second stage is 10 times more expensive then the average cost



A Practical Example

For testing, we generate multiple samples per instance

```
In [23]: data_ts = util.generate_costs(nsamples=150, nitems=nitems, seed=seed, sampling_seed=seed+1, noise=noise)
util.plot_df_cols(data_ts, figsize=figsize, title='Training Set', scatter=True)
```



By doing this, we get a more reliable evaluation of uncertainty



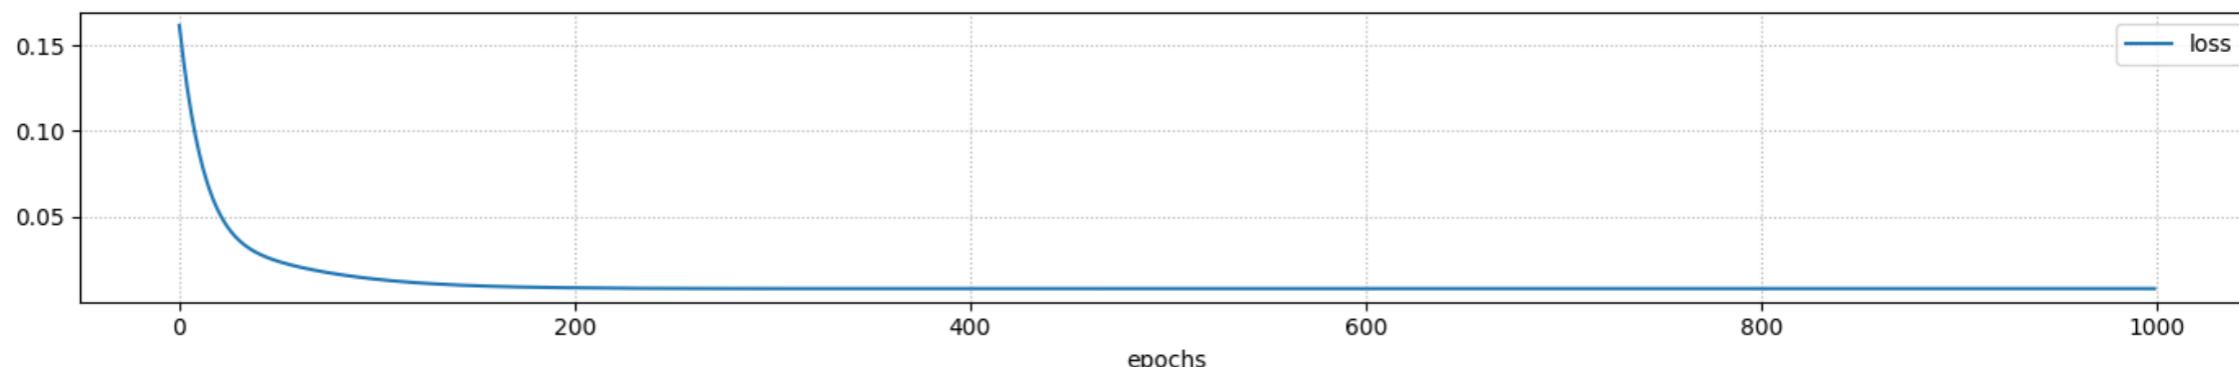
A PFL Approach

We start by training a prediction focused approach

```
In [24]: pfl_2s = util.build_nn_model(input_shape=1, output_shape=nitems, hidden=[], name='pfl_2s', output_activation='softmax')
@time history = util.train_nn_model(pfl_2s, data_tr.index.values, data_tr.values, epochs=1000, verbose=False)
util.plot_training_history(history, figsize=figsize_narrow, print_final_scores=False)
util.print_ml_metrics(pfl_2s, data_tr.index.values, data_tr.values, label='training')
util.print_ml_metrics(pfl_2s, data_ts.index.values, data_ts.values, label='training')
```

CPU times: user 8.21 s, sys: 352 ms, total: 8.56 s

Wall time: 6.83 s



R2: 0.80, MAE: 0.071, RMSE: 0.09 (training)

R2: 0.75, MAE: 0.072, RMSE: 0.09 (training)



This is as fast as the DFL approach, and can be used for warm-starting

Evaluating Two-Stage Approaches

Two-state stochastic approaches can be evaluated in two ways

We can compare them with **the best we could do**

- The cost difference is the proper regret
- Its computation requires solving a 2s-SOP with high accuracy
- ...Making it a very computationally expensive metric

We can compare them with the expected cost of **a clairvoyant approach**

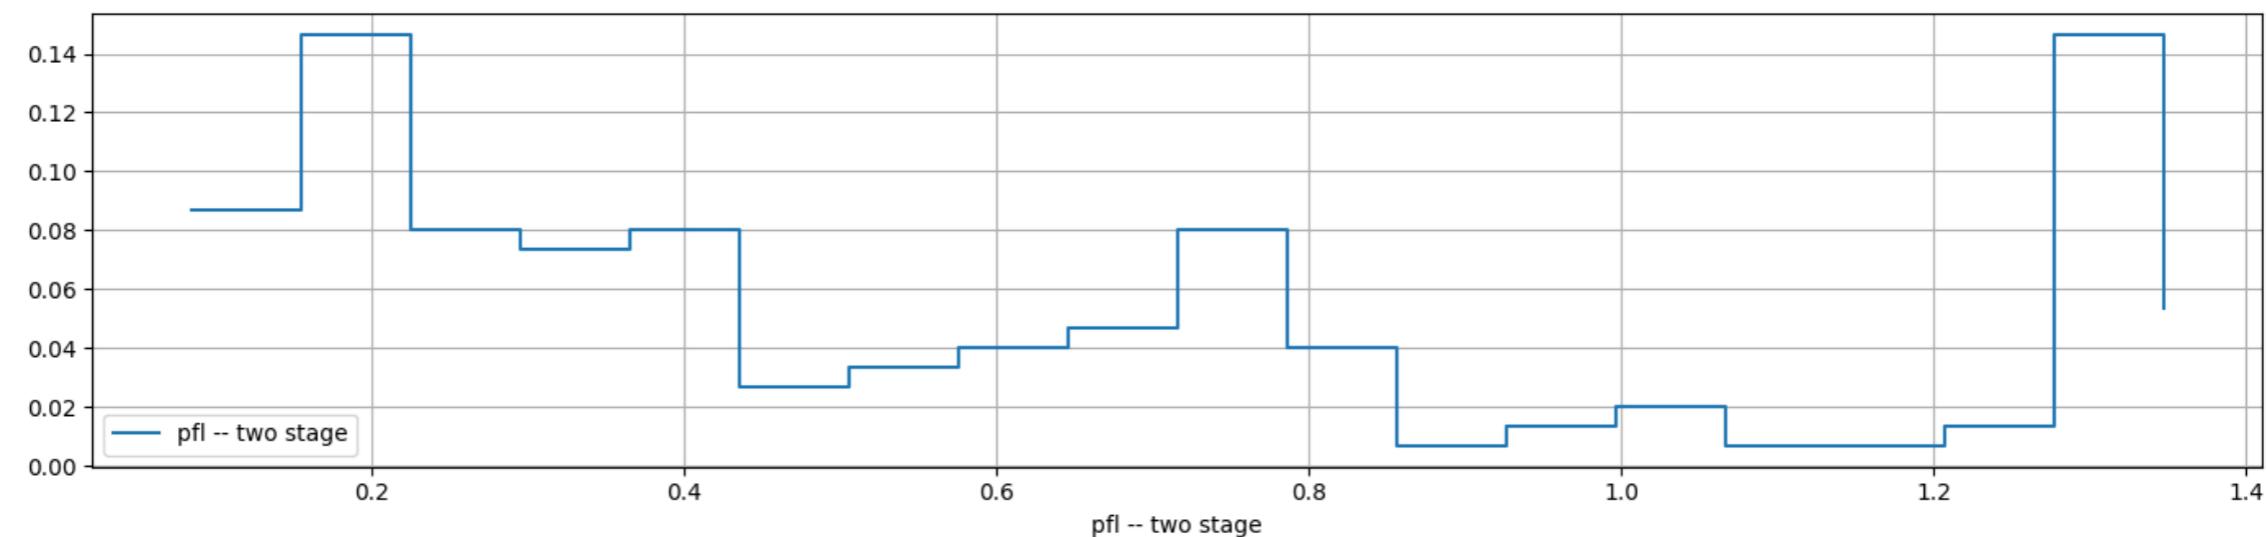
- The cost difference is called Expected Value of Perfect Information
- ...Or sometimes Post-hoc regret
- Its computation requires solving a 2s-SOP **with just a single scenario**
- ...So it's much faster, but only provide an upper bound on true regret



Evaluating the PFL Approach

Let's check the EVPF/Post-hoc regret for the PFL Approach

```
In [25]: pfl_2s_evpf = util.compute_evpf_2s(prb, pfl_2s, data_ts, tlim=10)
util.plot_histogram(pfl_2s_evpf, figsize=figsize, label='pfl -- two stage', print_mean=True)
```



Mean: 0.634 (pfl -- two stage)

This will be our baseline



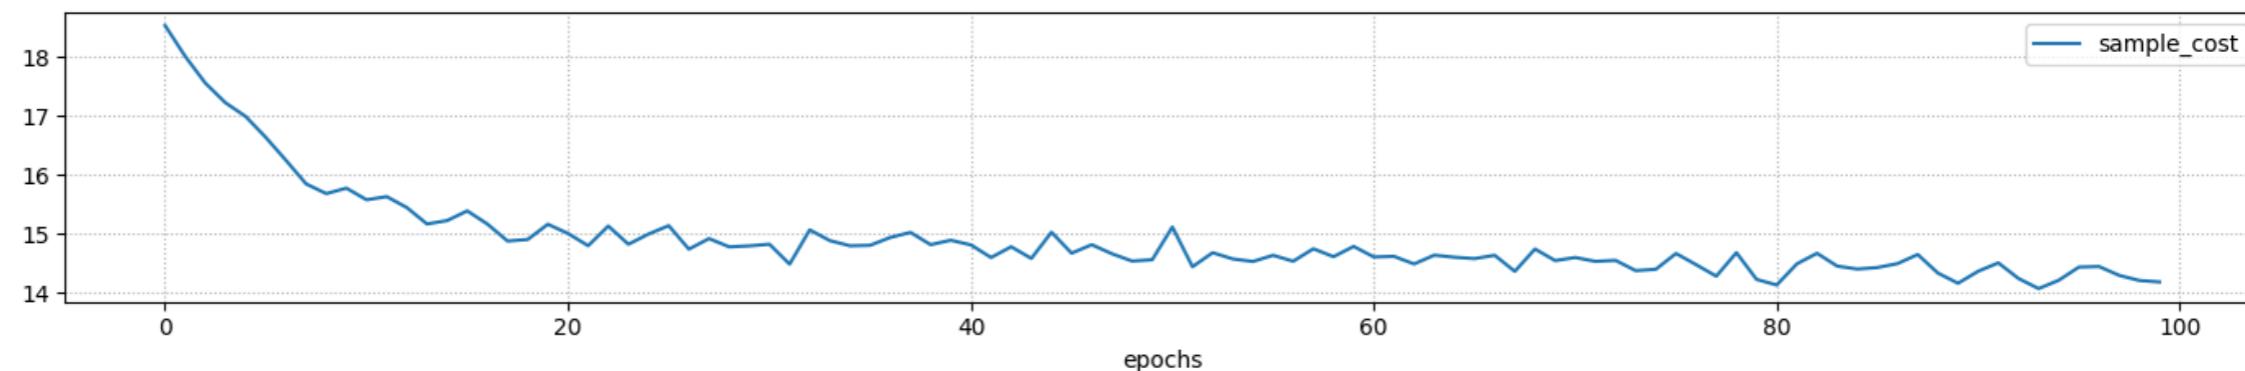
Training a DFL Approach

We traing a DFL with warm starting, but no solution cache

...Since the feasible space for the recourse actions is not fixed

```
In [30]: sfge_2s = util.build_dfl_ml_model(input_size=1, output_size=nitems, problem=prb, hidden=[], name='sample_cost')
@time history = util.train_dfl_model(sfge_2s, data_tr.index.values, data_tr.values, epochs=100,
util.plot_training_history(history, figsize=figsize_narrow, print_final_scores=False, excluded_r
util.print_ml_metrics(sfge_2s, data_tr.index.values, data_tr.values, label='training')
util.print_ml_metrics(sfge_2s, data_ts.index.values, data_ts.values, label='test')
```

CPU times: user 4min 41s, sys: 34.8 s, total: 5min 16s
Wall time: 5min 16s



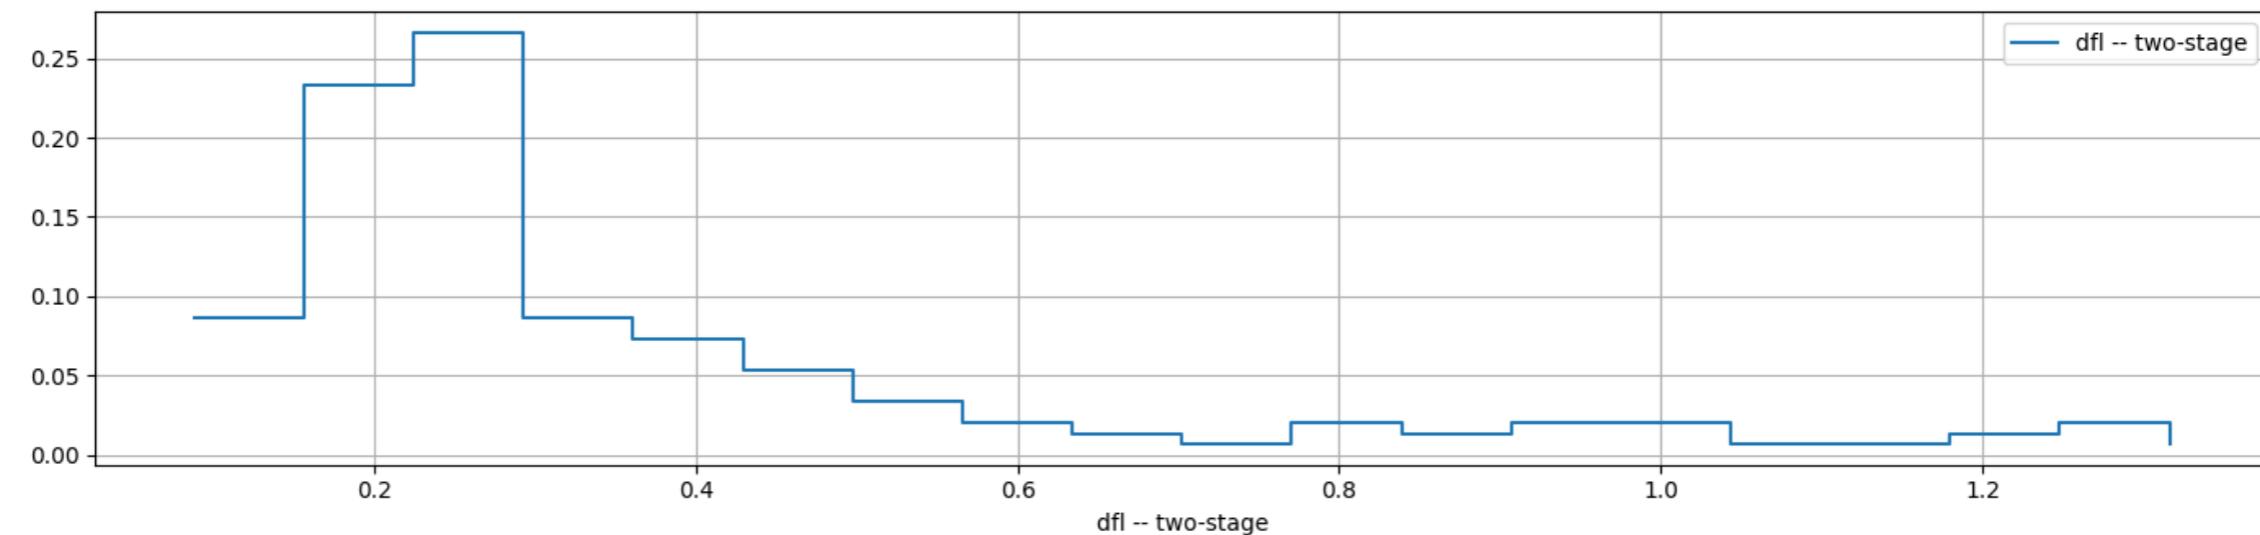
R2: 0.61, MAE: 0.095, RMSE: 0.12 (training)
R2: 0.66, MAE: 0.081, RMSE: 0.10 (test)



Evaluating the DFL Approach

We can now inspect the EVPF/Post-hoc regret for the DLF approach, as well

```
In [31]: sfge_2s_evpf = util.compute_evpf_2s(prb, sfge_2s, data_ts, tlim=10)
util.plot_histogram(sfge_2s_evpf, figsize=figsize, label='dfl -- two-stage', print_mean=True)
```



Mean: 0.382 (dfl -- two-stage)



A More In-depth Comparison

A more extensive experimentation will be found in [4]

The method has been tested on:

- Some "normal" DFL benchmarks
- Several two-stage stochastic problems

The baselines are represented by:

- Specialize methods (e.g. SPO, the one from [1]), when applicable
- A neuro-probabilistic model + a scenario based approach

Specialized method tend to work better

- ...But SFGE is much more versatile
- The best results are obtained on 2s-SOPs



A More In-depth Comparison

This is how the approach fares again the scenario based method

...On a problem somewhat similar to our supply planning one

